



Ardia Precision Health

Governed, non-diagnostic AI for **precision medicine** – spanning **molecular & genomic diagnostics**, the **2027 PAMA rate cliff**, **pulmonary & respiratory care**, and **elder care**.
Every answer cited to real policy – and withheld when it cannot be.

PRE-SEED · \$3-7M SAFE

NVIDIA INCEPTION MEMBER

PRECISION MEDICINE · PAMA · PULMONARY · ELDER CARE

ardiahealthlabs.com · founders@ardiahealthlabs.com · Live engine: ardiahealthlabs.com/studio

THE PROBLEM

Precision medicine finds the answer. The coverage decision buries it.

A tumour is sequenced. A matched therapy exists. The molecular test is denied under a Medicare coverage policy the ordering clinician never reads – and nobody appeals it. The same pattern repeats across pharmacogenomics, pulmonary and respiratory care, and the elder population.

23-31%

of molecular & genomic tests are **denied** by payers

~65%

of those denials are **never appealed** – silently written off

2x

the loss: **revenue** for the lab and a **patient** who misses the therapy

Why it persists: payers adjudicate against LCD/NCD and MoIDX/DEX policy with automated systems at algorithmic speed. Independent labs fight back with spreadsheets and billing staff who are not reimbursement-policy specialists. A structurally mismatched fight.

WHY NOW

The 2027 PAMA cliff makes every recovered dollar existential.

- **January 1, 2027:** CMS PAMA reimbursement cuts take effect, landing hardest on independent labs.
- Margins compress on exactly the molecular and toxicology panels these labs run.
- Denial recovery moves from back-office nice-to-have to a **survival line item**.
- Frontier reasoning models only recently became good enough to reason over policy – and cheap enough to run per claim.
- State AI-in-healthcare law (TX SB 1188, TRAIGA) now **rewards** auditable systems and punishes black boxes.

THE WINDOW

2027

PAMA cuts
go live

Labs are budgeting for that cliff **now**. The next few quarters are when they choose their denial-recovery strategy – and when a governed engine either gets designed in, or doesn't.

TARA – a governed reasoning engine between the evidence and the decision.

TARA reads the denial, checks it against the **actual governing Medicare coverage determination**, and drafts a citation-backed appeal a specialist approves – or withholds the answer entirely.



PROOF – CHECKABLE IN FIVE MINUTES

It's not a slide. It's live.

- Open **ardiahealthlabs.com/studio** and paste a synthetic denial.
- It returns **real Local Coverage Determinations – L35025 and L38045** – with working cms.gov links.
- A response that fails a gate is **withheld**, and the failed gate is named.
- Test suite public at **/test-results**. Framework at **/tara-framework**. Code on GitHub.

No demo call required. No sales gate. Every claim on this slide is verifiable by a stranger, today.

34/34

governance tests passing in CI

6

deterministic gates – a failure withholds the answer

L35025



a real LCD returned live, with a working cms.gov link

100%

of returned answers bound to a citation

One governed core. Four care domains.

A neuro-symbolic design: deterministic symbolic policy logic wrapped around a frontier reasoning model (Google Gemini today; provider-portable by design). The symbolic layer is what makes the output auditable.



Independent molecular & toxicology labs.

The profile

20-150 staff. Not owned by Quest or LabCorp. Running genomic, PGx and toxicology panels outside a hospital system. Billing Medicare and commercial payers under LCD/NCD and MoIDX/DEX.

Why them

Large enough that denied-claim volume is material. Small enough that manual appeals overwhelm a lean revenue-cycle team. Disproportionately exposed to the 2027 PAMA cliff.

ROLE	WHAT THEY OWN
Lab Director (MD/PhD)	Economic buyer · testing-platform authority
Revenue-Cycle Lead	Our champion – feels the denial pain daily
Technical Supervisor	Day-to-day workflow decisions
Lab Manager	Budget and staffing
Compliance Officer	The gate our governance is built to pass

Beachhead: one Dallas-Fort Worth lab. Named Phase-1 prospects: HealthTrackRx, Principle Health Systems, Filos Health Laboratory, Apollo.

We get paid only when the lab does.

15% contingency fee on
revenue we recover

- No subscription. No license. No implementation fee.
- Recover nothing → the lab owes **nothing**.
- Paid out of dollars already being written off – not a new budget line.
- Removes procurement friction for a buyer who cannot approve ROI on faith.

ILLUSTRATIVE UNIT ECONOMICS

Appealable denials / lab / yr	~\$1.0M
Recoverable via grounded appeal	~40%
Recovered for the lab	~\$400K
Ardia revenue (15%)	~\$60K ACV

Illustrative model from published denial rates – not measured results. Verifying per-lab denied-claim value is what our discovery interviews fund.

A multi-\$B precision-health layer, entered through one lab.

SOM • 12-24 MO

low \$M

A handful of DFW design-partner and early paid labs at ~\$60K ACV each. Small by design – sized to prove the model on real claims.

SAM • SERVICEABLE

\$100s M

~800-1,500 US independent molecular/tox labs, 20-150 staff, outside a chain or hospital. Count being verified, not asserted.

TAM • TOTAL

multi-\$B

US lab-reimbursement adjudication, then every adjacent decision the same core addresses – precision medicine & PGx, PAMA, pulmonary, elder care.

Sourced inputs

Molecular/genomic denial rates **23-31%**. Roughly **65%** never appealed. US clinical-lab testing is a **\$30B+** market; denial write-offs run to billions annually.

What we still owe you

A verified independent-lab count and a verified average denied-claim value per lab. We will not hand you a number we cannot defend – securing those is what our 15 discovery interviews are for.

COMPETITION

Everyone optimizes the repetitive middle. Nobody reasons over the policy tail.

CATEGORY	WHO	CITED TO POLICY	CAN WITHHOLD	BUILT FOR MOLECULAR TAIL
Generic RCM / denial mgmt	Waystar, R1 RCM, Availity, Candid	×	×	No – high-volume repetitive claims
Prior-auth automation AI	Cohere Health, Infinitus, Rhyme	partial	×	No – authorization side, not appeals
Lab RCM / LIS billing	XiFin, Telcor	static table	×	Partly – but policy is hand-maintained
General LLM copilots	Prompt-configured GPT wrappers	unenforced	×	No – grounding lives in a prompt
Ardia • TARA	Governed adjudication for labs	✓ enforced	✓ tested	Yes – that is the wedge

The honest read: the incumbents are strong, well funded, and good at what they do. But molecular and genomic testing is a small, policy-heavy tail for them – and general LLM copilots are precisely why RCM teams have grown sceptical of AI touching claims. We win where grounding must be **provable**, not promised.

Three layers – and we don't pretend the top one is the moat.

LAYER 1 • COMMODITIZING

Neuro-symbolic architecture

Genuinely novel in reimbursement adjudication today – but hybrid approaches will spread. This is the entry ticket, not the defense. We say so plainly.

LAYER 2 • COMPOUNDING

Executable governance

6 gates + 34 CI tests an auditor can run, and a failed gate withholds the answer. With compliance-officer buyers burned by AI demos, provable governance is a purchasing prerequisite.

LAYER 3 • DURABLE

Data & workflow lock-in

A continuously-maintained corpus of determinations mapped to reasoning logic; adjudicated outcomes from every lab sharpen the engine; deep switching costs once embedded in the appeal workflow.

Four method-claim provisionals scoped

Toxicology denial-recovery pipeline · MoIDX/DEX Z-code reasoning chain · PAMA-aware test-mix optimization · refusal-as-governance. Identified, not yet filed – counsel engaged at close.

Team advantage that can't be cloned fast

We engineered *inside* Cigna, UnitedHealth/Optum, Teladoc and ECFMG – the exact systems that generate these denials.

What's built, and what isn't.

✓ REAL TODAY

- Live engine returning **real LCDs L35025 / L38045**
- 6 deterministic gates · a failed gate withholds the answer
- **34/34 CI tests** passing, published publicly
- PHI de-identification running pre-inference
- **NVIDIA Inception** member · Google for Startups Cloud
- Named DFW Phase-1 pipeline (4 labs) · Delaware C-Corp, Dec 2025

× NOT YET – STATED PLAINLY

- **0 customers · 0 pilots · 0 signed LOIs**
- **\$0 revenue · \$0 raised** – bootstrapped
- Synthetic data only – no real-world validation
- De-identification catches structured IDs; **name detection in progress**
- PulmoIQ guideline retrieval corpus **still in development**
- HIPAA matrix self-graded **2 of 15** · no SOC 2, HITRUST, pen test, BAA

Why we publish the right-hand column: our buyers are compliance officers. The first claim they can't verify ends the relationship. We would rather be told our thesis is wrong early than be believed for the wrong reason.

TEAM

We built inside the systems that create these denials.



Rambabu (Ram) Vadlamudi

Founder • Principal AI / Data Architect

Sole technical builder of TARA. 15+ yrs enterprise engineering, 9+ in healthcare IT – Software Engineering Lead at **Cigna**, Senior SDE at **Teladoc**, Senior Software Developer at **ECFMG**, plus **UnitedHealth/Optum** and ALSAC/St. Jude.



Manasa Jampani

Co-Founder

10+ yrs payer and health-services operations across **UnitedHealth Group**, **Interwell Health** and **ECFMG** – the side that decides coverage and manages care. Owns operating and go-to-market judgment.



Paramesh Kurapati


Chief Executive Officer

Joined 2026 to lead strategy, go-to-market, partnerships and fundraising. A practicing veterinary doctor, he brings a clinical and scientific lens to how Ardia builds governed, non-diagnostic AI.

Board: Sireesha Mamillapalli, Ph.D. – Assistant Professor of Biochemistry & Physiology, Geisinger Commonwealth School of Medicine. Brings scientific rigor to our validation methodology. · **Actively recruiting:** a lab revenue-cycle advisor and a healthcare-AI regulatory advisor.

THE PLAN

One design partner unblocks everything.



USE OF FUNDS

- **One paid DFW lab pilot** in Q4 2026 under a DUA – real-data validation
- **15 customer-discovery interviews** – verify the two market inputs we owe
- **A claims-data engineer** who has shipped X12 835/837 in production
- Regulatory counsel opinion on SaMD scope + the security attestations labs require

THE BIGGEST RISK WE NAME

Data access, not technology. EDI 835/837 files are PHI, so validation requires a Limited Data Set + DUA with a risk-averse lab. That is a trust gate – and precisely why we de-identify before inference and make governance auditable in CI. Non-dilutive lane in parallel: **NSF SBIR Project Pitch.**

THE ASK

Raising a **\$3-7M** pre-seed on a SAFE.

No priced round, no cap set. Target close in 6-9 months. The capital takes a live, governed, testable engine to its first paid lab pilot – and turns synthetic proof into recovered revenue on real claims.

\$3-7M

Pre-seed SAFE · no cap

Q4 2026

First paid DFW pilot under a DUA

1 lab

The design partner that unblocks the roadmap



founders@ardiahealthlabs.com · +1 469 679 3334

Run the engine yourself: **ardiahealthlabs.com/studio**