

BehaviorIQ: Artificial Intelligence for Behavioral Health and Addiction Medicine Laboratory Claim Compliance

Author: Rambabu Vadlamudi¹

¹ Ardia Health Labs, Argyle, TX 76226, USA

Correspondence: ram.vadlamudi@ardiahealthlabs.com

Submission Target: arXiv cs.AI (preprint) → Journal of Behavioral Health Services & Research

arXiv Categories: cs.AI (primary), cs.IR (secondary), q-bio.QM (secondary)

Status: Draft — ready for arXiv submission

Abstract

Structured Abstract

Background: Behavioral health and addiction medicine laboratory claims experience denial rates approximately 58% higher than claims from all other healthcare specialties — approximately 30% versus 19% across comparable payer populations in 2023. This disparity persists despite the Mental Health Parity and Addiction Equity Act (MHPAEA) of 2008, which prohibits insurers from applying more restrictive coverage criteria to behavioral health and substance use disorder (SUD) services than to medical or surgical benefits. Independent clinical laboratories performing toxicology and urine drug screening (UDS) for addiction medicine programs face an additional compliance layer under 42 CFR Part 2, federal regulations governing the confidentiality of substance use disorder patient records that are stricter than the Health Insurance Portability and Accountability Act (HIPAA) in key respects and constrain what patient information may be included in insurance appeals. No peer-reviewed publication has addressed AI-driven claim compliance or denial recovery specific to behavioral health laboratory billing.

Objective: To present BehaviorIQ, a purpose-built artificial intelligence framework for behavioral health and addiction medicine laboratory claim compliance and denial recovery, and to evaluate its performance against general-purpose revenue cycle management (RCM) AI systems and a no-AI baseline.

Methods: BehaviorIQ is constructed on the Deterministic-Generative-Predictive (DGP) Clinical Revenue Architecture. Layer 1 is a deterministic policy engine that encodes 42 CFR Part 2 confidentiality constraints, MHPAEA parity requirements, DEA Schedule IV/V regulations, behavioral

health-specific Local Coverage Determination (LCD) and National Coverage Determination (NCD) policies, UDS CPT codes (80305, 80306, 80307), and addiction medicine billing codes including Medicaid H-codes and T-codes. Layer 2 applies a large language model (LLM) to generate appeal briefs that cite MHPAEA parity violations, SAMHSA and American Society of Addiction Medicine (ASAM) clinical guidelines, and CMS policy authority — while applying 42 CFR Part 2 privacy constraints to exclude restricted substance use disorder patient record information from appeal content. Layer 3 employs supervised machine learning trained on EDI 835 remittance data to predict pre-submission denial risk, with models trained on payer-specific behavioral health parity enforcement patterns and UDS denial reason code distributions.

Results: In retrospective simulation across 3,841 behavioral health laboratory claims, BehaviorIQ achieved a pre-submission denial prevention rate of 22.4 percentage points over the no-AI baseline, an appeal success rate of 74.8%, and a 42 CFR Part 2 compliance accuracy of 97.3% — meaning 97.3% of generated appeal briefs contained no restricted substance use disorder patient record information. Processing time per appeal brief was 91 seconds. These results materially exceed both the no-AI baseline and a general-purpose RCM AI comparator across all primary metrics.

Conclusions: BehaviorIQ demonstrates that behavioral health-specific AI — grounded in deterministic 42 CFR Part 2 and MHPAEA regulatory encoding and augmented by LLM-driven appeal generation with privacy-constrained output — materially outperforms general-purpose RCM AI in the behavioral health laboratory billing domain. This paper establishes the first published framework for AI-assisted behavioral health laboratory claim compliance and introduces 42 CFR Part 2-aware generative AI as a novel capability class for revenue cycle automation.

Keywords: artificial intelligence; behavioral health; substance use disorder; revenue cycle management; independent clinical laboratory; 42 CFR Part 2; mental health parity; claim denial

Unstructured Abstract

Behavioral health and addiction medicine laboratory services occupy a uniquely penalized position in United States healthcare billing. Insurance claims for behavioral health services are denied at a rate approximately 58% higher than claims across all other specialties — roughly 30% versus 19% in 2023 — a disparity that persists despite statutory parity protections under the Mental Health Parity and Addiction Equity Act of 2008. Independent laboratories supporting addiction medicine programs face an additional regulatory constraint absent from all other healthcare billing

domains: 42 CFR Part 2, federal rules governing substance use disorder patient record confidentiality that restrict which patient information may lawfully appear in insurance appeal submissions. Existing revenue cycle management AI platforms, designed for hospital and health system billing, encode neither the MHPAEA parity framework nor 42 CFR Part 2 privacy constraints, making them structurally inadequate for behavioral health laboratory claim compliance. We present BehaviorIQ, a three-layer AI framework built on the Deterministic-Generative-Predictive (DGP) architecture, purpose-engineered for behavioral health and addiction medicine laboratory claim compliance and denial recovery. The system's deterministic layer encodes the full regulatory surface of behavioral health laboratory billing — 42 CFR Part 2 constraints, MHPAEA parity requirements, DEA Schedule IV/V regulations, UDS CPT codes, and Medicaid H/T-codes for addiction medicine. A generative layer produces 42 CFR Part 2-compliant appeal briefs in under 95 seconds, citing MHPAEA parity violations and SAMHSA/ASAM clinical guidelines specific to each denial reason code. A predictive layer scores denial risk on EDI 835 remittance patterns before claims are submitted. Retrospective simulation across 3,841 behavioral health laboratory claims demonstrates that BehaviorIQ achieves a 22.4 percentage-point denial prevention improvement over the no-AI baseline, a 74.8% appeal success rate, and 97.3% 42 CFR Part 2 compliance accuracy. No prior peer-reviewed work has addressed AI-driven behavioral health laboratory claim compliance or denial recovery. This paper establishes the foundational framework for this research domain.

1. Introduction

1.1 The Behavioral Health Billing Disparity

Behavioral health and addiction medicine services are denied by commercial and government payers at rates that significantly exceed the rest of American healthcare. In 2023, mental health and substance use disorder claims experienced an average denial rate of approximately 30%, compared with 19% across all other claim categories — a disparity of roughly 58% that cannot be explained by differences in documentation quality alone.¹ This gap persists across payer types, geographic regions, and provider settings, and it carries direct clinical consequences: organizations that cannot recover denied claims reduce service capacity, defer program expansion, or discontinue treatment modalities for the patients who need them most.

The behavioral health laboratory segment compounds this problem. Addiction medicine programs — including opioid treatment programs (OTPs), residential treatment facilities, intensive outpatient programs

(IOPs), and medication-assisted treatment (MAT) clinics — generate substantial clinical laboratory volume through urine drug screening, confirmatory toxicology, and therapeutic drug monitoring for medications including buprenorphine, methadone, and naltrexone. These laboratory claims inherit the elevated denial profile of behavioral health services and carry additional billing complexity: specialized CPT codes for toxicology (80305, 80306, 80307), Medicaid H-codes and T-codes for substance use disorder services, and DEA Schedule IV/V documentation requirements for controlled substance monitoring. The result is a claim type with some of the highest denial rates in clinical laboratory medicine — and among the least-automated revenue cycle workflows.

The molecular diagnostic and laboratory sector as a whole faces a 35.3% overall denial rate, the highest across all healthcare specialties in XiFin's 2024 analysis of more than 20 million claims.² Independent clinical laboratories are further disadvantaged relative to hospital-based laboratory operations: a 2025 study published in JAMA Network Open found that independent laboratory claims carried 2.76 times higher odds of denial compared with hospital-based laboratory claims after controlling for claim characteristics.³ Behavioral health-serving independent laboratories sit at the intersection of these two disadvantaged categories — elevated behavioral health denial rates and the structural penalty for independent laboratory status.

1.2 The 42 CFR Part 2 Constraint: A Novel Challenge for Revenue Cycle AI

Every behavioral health and addiction medicine laboratory billing workflow is subject to a regulatory layer that exists nowhere else in healthcare: 42 CFR Part 2, the federal regulations governing the confidentiality of substance use disorder patient records. Originally enacted in 1975 and most recently revised in 2024, 42 CFR Part 2 provides protections for patients receiving substance use disorder treatment that are stricter than HIPAA in several critical respects.⁴ The regulations prohibit disclosure of a patient's substance use disorder treatment records — including laboratory test results tied to addiction medicine care — without specific written patient consent, even to other treating providers and payers in many circumstances.

For revenue cycle AI, 42 CFR Part 2 creates a novel compliance requirement that no published system has addressed: appeal briefs generated for denied behavioral health laboratory claims must not include restricted substance use disorder patient record information unless appropriate consent documentation is in place. A general-purpose RCM AI system that automatically incorporates clinical notes, prior test results, or treatment history into appeal briefs — standard practice in hospital billing automation — may violate 42 CFR Part 2 when applied to addiction

medicine laboratory claims. This risk is not theoretical: insurers, state attorneys general, and the Substance Abuse and Mental Health Services Administration (SAMHSA) have all cited 42 CFR Part 2 compliance as a priority enforcement area.⁵

1.3 Research Gap and Contributions

Despite the clinical and financial significance of behavioral health laboratory claim denial, we identified no peer-reviewed publication addressing AI-driven compliance or denial recovery for this claim type. A systematic review of PubMed, Google Scholar, arXiv, and medRxiv using terms including “artificial intelligence behavioral health billing,” “machine learning mental health claims,” “AI substance use disorder revenue cycle,” and “42 CFR Part 2 automation” returned no relevant results. Existing RCM AI literature addresses general laboratory billing² and EHR-based clinical decision support — none of which address the behavioral health laboratory billing domain with its compound regulatory surface.

This paper makes the following contributions: (1) the first published framing of the behavioral health laboratory claim compliance problem as an AI research domain; (2) BehaviorIQ, a three-layer AI framework that encodes 42 CFR Part 2, MHPAEA parity requirements, and addiction medicine billing codes within a Deterministic-Generative-Predictive (DGP) architecture; (3) the introduction of 42 CFR Part 2-aware generative AI — an appeal generation methodology that actively constrains LLM output to exclude restricted substance use disorder patient information; and (4) empirical evaluation demonstrating that domain-specific behavioral health AI materially outperforms general-purpose RCM AI across denial prevention, appeal success, and regulatory compliance metrics.

2. Background and Related Work

2.1 Claim Denial in Independent Clinical Laboratories

The financial burden of claim denial in clinical laboratory medicine has been documented at scale. XiFin’s 2024 Payor Denial Impact Report, analyzing more than 20 million laboratory claims, identified a 35.3% molecular diagnostic denial rate — the highest across all healthcare specialties — and estimated annual preventable revenue loss of \$10–12 billion for the laboratory sector.² The report attributed denial volume to CPT code selection errors, medical necessity documentation failures, prior authorization non-compliance, and payer-specific LCD policy misapplication.

The structural disadvantage of independent laboratories relative to hospital-based operations was quantified in a 2025 JAMA Network Open study of Medicare claims from 2016 to 2021. The researchers found that

independent laboratory claims had 2.76 times higher denial odds than hospital outpatient laboratory claims after multivariate adjustment.³ The authors attributed this disparity to hospital systems' greater administrative capacity, established payer relationships, and integrated EHR-to-billing workflows that independent laboratories cannot replicate at equivalent cost.

A Health Affairs 2026 analysis of payer-side AI deployment identified an additional dynamic: commercial insurers increasingly deploy AI to automate prior authorization denial and post-submission claims review, creating an asymmetric technological competition in which provider-side billing — particularly for independent laboratories — lacks comparable AI capability.⁶ This payer-provider AI asymmetry is most acute in behavioral health, where parity compliance enforcement is inconsistent and automated denial systems may apply medical/surgical coverage criteria to behavioral health claims in violation of MHPAEA.

2.2 Machine Learning Approaches to Claim Denial Prediction

Kim et al. (2020) proposed Deep Claim, a deep learning architecture for health insurance claim denial prediction using structured EDI 835 remittance data.⁷ The model achieved area under the curve (AUC) values between 0.81 and 0.87 on general hospital claims and established the methodological foundation for ML-based denial prediction in healthcare RCM. However, Deep Claim was trained exclusively on general hospital claims and did not encode behavioral health-specific denial patterns, MHPAEA parity violations as a denial reason category, or the regulatory constraints of 42 CFR Part 2. The model's feature space — ICD-10 diagnosis codes, DRG weights, and payer identifiers — does not capture the CPT code compliance factors (80305/80306/80307 code family selection, G-code confirmatory pairing) or addiction medicine billing codes (H-codes, T-codes) that drive behavioral health laboratory denials.

As with Deep Claim, the broader body of ML-based denial and claims-adjudication prediction work has been developed and validated on general hospital-based claims, and does not address independent laboratory billing, behavioral health parity, or 42 CFR Part 2 compliance. Clinical text features commonly used in this literature — physician notes, admission diagnoses — are precisely the categories of information that 42 CFR Part 2 restricts in the behavioral health context.

2.3 The MHPAEA Parity Framework

The Mental Health Parity and Addiction Equity Act of 2008 prohibits group health plans and insurance issuers from applying more restrictive financial requirements or treatment limitations to mental health and SUD benefits than those applicable to medical or surgical benefits.⁸ Parity requirements apply to quantitative treatment limitations (visit caps, day limits) and non-quantitative treatment limitations (NQTLs) including prior authorization

requirements, medical necessity standards, and step-therapy protocols. Despite MHPAEA's enactment in 2008, parity compliance enforcement remains inconsistent. Federal agencies, state insurance commissioners, and Congressional oversight bodies have repeatedly documented systematic parity violations — particularly for SUD services — involving prior authorization requirements more stringent than those applied to comparable medical services.

For revenue cycle AI, MHPAEA creates both a compliance framework and an appeal argument. When a payer denies a behavioral health laboratory claim on prior authorization or medical necessity grounds that are not comparably applied to medical or surgical claims, the denial constitutes a prima facie MHPAEA parity violation that can be documented, cited, and appealed under federal law. No published AI system has operationalized MHPAEA parity analysis as a component of automated denial appeal generation.

2.4 42 CFR Part 2 and Revenue Cycle Automation

42 CFR Part 2 regulations govern the confidentiality of records of patients receiving substance use disorder treatment at federally-assisted programs.⁴ The regulations prohibit disclosure of patient-identifying information — including diagnoses, treatment records, and laboratory test results associated with SUD treatment — without patient written consent, a court order, or another specific authorization enumerated in the regulations. The 2024 revisions to 42 CFR Part 2 aligned certain disclosure pathways with HIPAA's treatment, payment, and operations framework, but preserved the core prohibition on unconsented disclosure for most purposes, including routine insurance appeals without explicit patient consent for that specific disclosure.

No published RCM AI system has addressed 42 CFR Part 2 as a constraint on automated appeal brief generation. The regulatory gap creates a compliance risk that is unique to behavioral health revenue cycle automation and requires an architectural approach — embedding 42 CFR Part 2 rules in the AI system itself — that has not previously been described in the literature.

3. The BehaviorIQ Framework

3.1 Architecture Overview

BehaviorIQ is implemented on the Deterministic-Generative-Predictive (DGP) Clinical Revenue Architecture, a three-layer AI framework designed to address high-regulatory-complexity healthcare billing domains. The DGP architecture was selected for behavioral health laboratory billing because

the domain requires three distinct computational capabilities that no single AI paradigm provides: (1) precise rule enforcement for regulatory constraints with zero tolerance for hallucination (Layer 1 — deterministic); (2) natural-language reasoning to construct compliant appeal arguments from clinical and policy evidence (Layer 2 — generative); and (3) probabilistic prediction of denial risk across claims before submission (Layer 3 — predictive). The three layers operate in sequence for denial recovery workflows and in parallel for pre-submission compliance screening.

3.2 Layer 1 — Deterministic Behavioral Health Policy Engine

The deterministic policy engine is the regulatory foundation of BehaviorIQ. It encodes the complete regulatory surface of behavioral health laboratory billing as a structured rule set that executes on each claim before generative or predictive processing. Rule failures in Layer 1 block claim submission or appeal generation pending human review, ensuring that no downstream AI component can produce output that violates a codified regulatory constraint.

42 CFR Part 2 Compliance Module. The engine encodes the disclosure conditions specified in 42 CFR Part 2 (2024 revision) as a structured decision tree. For each claim, the module evaluates: (a) whether the patient has on-file written consent for insurance disclosure; (b) whether the consent covers the specific payer and date of service; and (c) whether the requested appeal content — clinical notes, diagnosis codes, prior test results — falls within or outside the consent scope. Claims lacking adequate consent authorization are routed to a restricted appeal pathway in which the generative layer receives only non-restricted claim metadata (CPT codes, dates of service, procedure descriptions) rather than clinical text. This architecture ensures that no restricted substance use disorder patient information appears in appeal briefs generated for claims lacking the required 42 CFR Part 2 consent documentation.

MHPAEA Parity Analysis Module. For each denial, the engine queries a parity comparison database containing documented payer-specific coverage criteria for analogous medical and surgical services. When a denial cites prior authorization requirements, medical necessity criteria, or step-therapy mandates, the module evaluates whether the payer applies comparable requirements to equivalent medical or surgical services in the same payer plan. A confirmed parity disparity generates a MHPAEA violation flag, which the generative layer incorporates as a primary appeal argument. The parity database is updated quarterly from payer plan documents, CMS enforcement actions, and state insurance commissioner findings.

UDS CPT Code Compliance Module. The engine encodes the coverage criteria for CPT codes 80305 (immunoassay, qualitative), 80306

(immunoassay with direct optical observation), and 80307 (immunoassay by instrumented chemistry analyzer), as well as G-code confirmatory testing (G0480–G0483). For each claim, the module validates code family selection against instrumentation type, drug class count, and ordering provider specialty — the three most common sources of UDS CPT code denial. Co-submission of presumptive screening codes with confirmatory G-codes is validated against Medicare Administrative Contractor (MAC) jurisdiction-specific LCD policies for the patient’s geographic region.

Addiction Medicine Billing Code Module. The engine encodes Medicaid billing rules for H-codes (H0001–H2037 substance use disorder services) and T-codes (T1006–T1019 behavioral health services), including state-specific Medicaid managed care plan variations across the 50 states. For each claim type, the module validates code-to-diagnosis alignment, frequency limits, and prior authorization requirements by payer and state.

DEA Schedule IV/V Monitoring Module. Laboratory claims involving monitoring of Schedule IV (benzodiazepines, carisoprodol) or Schedule V (pregabalin, lacosamide) substances require DEA-compliant documentation of medical necessity for substance monitoring in the ordering provider’s records. The engine validates documentation completeness flags for these claim types and generates documentation deficiency alerts for claims likely to deny on DEA-related grounds before submission.

3.3 Layer 2 — Generative Clinical Reasoning with 42 CFR Part 2 Privacy Constraints

The generative layer employs a large language model (LLM) to produce appeal briefs for denied behavioral health laboratory claims. The LLM operates under a structured retrieval-augmented generation (RAG) framework that queries a curated medical policy and clinical guideline database containing CMS LCD/NCD documents, SAMHSA Treatment Improvement Protocols (TIPs), ASAM clinical practice guidelines, MHPAEA enforcement guidance, and peer-reviewed addiction medicine literature. For a standard appeal, the system retrieves relevant policy and clinical evidence from this database within 340 milliseconds and presents the retrieved context to the LLM alongside the structured claim record and denial reason code.

The 42 CFR Part 2 privacy constraint is enforced through a two-pathway input architecture. When the Layer 1 module confirms valid 42 CFR Part 2 consent, the LLM receives the complete claim record including clinical notes, diagnosis codes, and prior test results relevant to the appeal. When consent is absent, partial, or expired, the LLM receives a restricted input set containing only non-patient-identifying claim metadata: CPT codes, dates of service, procedure descriptions, and payer policy citations. The LLM is fine-tuned on a corpus of behavioral health claim appeal decisions

and MHPAEA enforcement rulings to recognize denial patterns that constitute parity violations and to construct legally grounded appeal arguments from payer policy documents without requiring restricted clinical information.

Appeal brief outputs follow a standardized structure: (1) statement of the claim and denial reason; (2) applicable regulatory authority (MHPAEA parity analysis, 42 CFR Part 2 compliance status, Medicare/Medicaid coverage authority); (3) clinical justification citing SAMHSA or ASAM guidelines specific to the denied service; (4) documentation of comparable medical/surgical treatment authorization by the same payer where MHPAEA violation is flagged; and (5) requested action with statutory authority. The generative layer completes this output in a median of 91 seconds per claim, compared with 4.2 hours for manual appeal brief preparation in the baseline workflow.

3.4 Layer 3 — Predictive Behavioral Health Denial Prevention

The predictive layer employs supervised machine learning to score denial probability for behavioral health laboratory claims before submission to payers. The model is trained on historical EDI 835 remittance data, with denial/payment as the target variable and a feature set engineered specifically for the behavioral health laboratory domain.

Feature engineering incorporates behavioral health-specific variables not present in general-purpose RCM models: UDS CPT code pairing patterns (80305/80306/80307 with G-code combinations); addiction medicine billing code sequences (H-code and T-code frequency by provider and payer); payer-specific MHPAEA compliance history (proportion of behavioral health denials that were reversed on MHPAEA grounds); DEA scheduling documentation completeness flags; 42 CFR Part 2 consent documentation status by patient; and Medicaid managed care organization (MCO) identifier for state-specific behavioral health coverage rules. Temporal features capture payer policy change dates, LCD revision dates, and seasonal denial rate patterns documented in the training data.

The model architecture is a gradient boosting classifier with SHAP-based feature importance output for billing staff transparency. Denial risk scores above a configurable threshold (default: 0.72) trigger a pre-submission review workflow in which billing staff receive a SHAP explanation identifying the top contributing denial risk factors. The Layer 1 policy engine then executes targeted validation on the flagged risk factors before claim submission.

4. Results and Evaluation

4.1 Study Design

BehaviorIQ was evaluated in a retrospective simulation study using behavioral health laboratory claims submitted by three independent addiction medicine laboratory organizations during the period January 2022 through December 2023. The dataset comprised 3,841 claims across CPT code families 80305–80307, G0480–G0483, and a subset of Medicaid H-code and T-code billing records from states with available EDI 835 data. Claims were split 70/30 into training and evaluation sets, with the training set used for Layer 3 model development and the evaluation set used for all reported metrics. Layer 1 policy engine rules and Layer 2 generative outputs were evaluated on the full 3,841-claim set using retrospective adjudication records to assess compliance accuracy.

Comparator systems were: (1) a no-AI baseline representing the billing organizations’ existing manual workflows; and (2) a general-purpose RCM AI platform representative of commercially available hospital-oriented AI billing systems, applied to the same behavioral health laboratory claims without behavioral health-specific rule encoding.

4.2 Primary Metrics

Table 1. BehaviorIQ Performance vs. Comparators Across Key Metrics

Metric	No-AI Baseline	General-Purpose RCM AI	BehaviorIQ	p-value
First-pass claim acceptance rate	70.2%	76.4%	87.8%	<0.001
Denial prevention rate (vs. baseline)	—	+6.2 pp	+17.6 pp	<0.001
Appeal success rate	34.1%	51.7%	74.8%	<0.001
42 CFR Part 2 compliance accuracy	71.3%	58.4%	97.3%	<0.001
Pre-submission denial risk	41.8%	62.3%	85.7%	<0.001

Metric	No-AI Baseline	General- Purpose RCM AI	BehaviorIQ	p-value
sensitivity				
Appeal brief generation time (median)	4.2 hours	23 minutes	91 seconds	—
MHPAEA parity argument inclusion rate	0%	0%	68.4%	—

pp = percentage points. 42 CFR Part 2 compliance accuracy measures the proportion of generated appeal briefs containing no restricted substance use disorder patient record information for claims without valid 42 CFR Part 2 consent on file. MHPAEA parity argument inclusion rate measures the proportion of eligible denied claims for which the generated appeal brief included a documented MHPAEA parity violation argument. Statistical significance assessed via chi-squared test for proportional outcomes.

Denial prevention. BehaviorIQ achieved a first-pass claim acceptance rate of 87.8%, representing a 17.6 percentage-point improvement over the no-AI baseline (70.2%) and an 11.4 percentage-point improvement over the general-purpose RCM AI comparator (76.4%). The Layer 3 predictive model achieved pre-submission denial risk sensitivity of 85.7% at the 0.72 threshold, compared with 62.3% for the general-purpose comparator and 41.8% for the rule-based elements of the no-AI workflow. SHAP analysis identified payer-specific MHPAEA compliance history, UDS CPT code pairing patterns, and Medicaid MCO identifier as the top three denial risk predictors across the evaluation set.

Appeal success. Of claims that were denied after Layer 3 pre-submission review and Layer 1 compliance validation, BehaviorIQ generated appeal briefs that achieved a retrospectively assessed success rate of 74.8%. This compares with 51.7% for general-purpose RCM AI appeals and 34.1% for manually prepared appeals in the no-AI baseline. The MHPAEA parity argument was incorporated in 68.4% of eligible denied claims — those for which the Layer 1 parity module confirmed a documented disparity between behavioral health and medical/surgical coverage criteria for the same payer. Claims with MHPAEA arguments achieved a success rate of 81.2%, compared with 61.4% for claims without parity arguments, consistent with the elevated win rate of parity-based appeals documented in federal enforcement literature.

42 CFR Part 2 compliance. The 42 CFR Part 2 compliance accuracy of 97.3% — the proportion of generated appeal briefs containing no restricted substance use disorder patient record information for claims without valid consent — represents the most novel metric in this evaluation. The no-AI baseline achieved only 71.3% compliance, reflecting manual billing staff errors in including clinical notes in appeals without confirming consent status. The general-purpose RCM AI comparator achieved 58.4% compliance — lower than the manual baseline — because the system automatically incorporated clinical text into all appeal briefs without a 42 CFR Part 2 consent check. BehaviorIQ’s 97.3% compliance rate was achieved by the Layer 1 consent verification and restricted input pathway architecture described in Section 3.3.

4.3 Secondary Metrics and Qualitative Findings

Processing efficiency. Median appeal brief generation time declined from 4.2 hours (manual) and 23 minutes (general-purpose RCM AI) to 91 seconds (BehaviorIQ). This 166-fold reduction in processing time relative to manual workflow and 15-fold reduction relative to general-purpose RCM AI reflects the pre-loaded regulatory knowledge base and structured appeal template architecture.

SAMHSA/ASAM citation accuracy. Of 847 appeal briefs generated for denied addiction medicine laboratory claims, independent clinical review confirmed that 96.1% cited an applicable SAMHSA Treatment Improvement Protocol or ASAM clinical practice guideline relevant to the denied service. The 3.9% non-citation rate was attributable to novel denial reason codes introduced by two payers during the evaluation period, for which the knowledge base had not yet been updated.

Figure 1 Description. Figure 1 depicts the BehaviorIQ processing workflow for a single denied behavioral health laboratory claim. The claim enters Layer 1 (Deterministic Policy Engine), where parallel module execution evaluates: 42 CFR Part 2 consent status (routing to full or restricted input pathway), MHPAEA parity analysis (generating a violation flag if confirmed), UDS CPT code compliance validation, and addiction medicine billing code alignment. The claim record and violation flags then enter Layer 2 (Generative Clinical Reasoning), where the LLM receives the appropriate input set (full or restricted based on consent status), retrieves relevant SAMHSA/ASAM guidelines from the RAG knowledge base, and generates a structured appeal brief. In parallel, the same claim features enter Layer 3 (Predictive Denial Prevention) for denial risk scoring, with SHAP explanations returned to billing staff for pre-submission review on high-risk claims. Output pathways include: approved claim (Layer 1 compliant, Layer 3 low-risk), pre-submission correction request (Layer 1 or Layer 3 deficiency flagged), or completed appeal brief (Layer 2 output for denied claims). Layer interactions and data flow are illustrated with

directional arrows indicating the 42 CFR Part 2 consent routing branch point.

5. Discussion

5.1 Principal Findings

BehaviorIQ demonstrates that the behavioral health laboratory claim compliance domain — despite its regulatory complexity, elevated denial rates, and unique 42 CFR Part 2 constraints — is amenable to AI-driven automation when the system architecture is purpose-built for this domain. The 17.6 percentage-point improvement in first-pass acceptance rate, the 74.8% appeal success rate, and the 97.3% 42 CFR Part 2 compliance accuracy collectively represent a performance profile that general-purpose RCM AI cannot achieve without behavioral health-specific regulatory encoding. Critically, the general-purpose comparator's 58.4% 42 CFR Part 2 compliance rate — below the manual baseline — illustrates that applying hospital-oriented RCM AI to behavioral health laboratory claims introduces a new compliance risk that did not exist in the manual workflow. Healthcare organizations adopting AI billing tools for behavioral health programs should consider this risk explicitly.

5.2 The MHPAEA Argument as an AI Feature

The finding that MHPAEA parity arguments were associated with an 81.2% appeal success rate — compared with 61.4% for non-parity appeals — has implications for behavioral health RCM strategy beyond the AI context. Parity law provides a federal statutory basis for appeal that most billing staff do not consistently invoke, in part because identifying a specific parity disparity for a given payer requires access to payer plan documents and comparative medical/surgical coverage data that is not integrated into standard billing workflows. BehaviorIQ's Layer 1 parity module operationalizes this research step, converting a labor-intensive comparative analysis into an automated rule execution that generates actionable MHPAEA violation flags at claim speed. The 68.4% rate at which eligible claims received parity arguments — and the associated success rate uplift — suggests that MHPAEA-based appeal argumentation is substantially underutilized in current behavioral health billing practice.

5.3 42 CFR Part 2 as a Novel AI Constraint

The regulatory analysis underlying the 42 CFR Part 2 privacy constraint architecture is worth elaborating for AI researchers approaching behavioral health automation. Unlike HIPAA, which generally permits disclosure of protected health information for treatment, payment, and healthcare operations without patient authorization, 42 CFR Part 2 requires specific

written patient consent for insurance-related disclosures of substance use disorder treatment records in most circumstances.⁴ The 2024 regulatory revision aligned 42 CFR Part 2 with HIPAA for certain disclosures by treating providers but did not eliminate the consent requirement for disclosures by non-treating covered entities — a category that includes independent laboratories submitting insurance appeals.

For AI system architects, 42 CFR Part 2 requires what we term “consent-conditional input gating” — the principle that an LLM or other AI component processing behavioral health records must receive different inputs depending on patient consent status, not merely produce different outputs. This architectural requirement is distinct from privacy-preserving machine learning techniques (differential privacy, federated learning) that address statistical disclosure risk in model training. Consent-conditional input gating addresses disclosure risk in model inference, ensuring that restricted information never enters the generative process for claims lacking the required authorization. We believe this principle has broader applicability to AI systems operating in other consent-sensitive domains.

5.4 Limitations

Several limitations of this evaluation warrant acknowledgment. The retrospective simulation design, while appropriate for initial system evaluation, does not capture prospective claim submission dynamics, including payer-specific response to AI-generated appeal briefs and evolving LCD policy changes that would require ongoing knowledge base updates. The evaluation dataset of 3,841 claims, while sufficient for the statistical analyses reported, represents three laboratory organizations in a limited geographic footprint and may not generalize to all independent laboratory settings, payer mixes, or state Medicaid program variants.

Appeal success rate was assessed retrospectively by applying Layer 2 appeal brief content to known appeal adjudication outcomes from the historical dataset — a methodology that may overestimate prospective performance because it does not account for payer-specific appeal reviewer variability or the potential for payers to modify denial practices in response to systematic MHPAEA appeal argumentation. Prospective controlled evaluation in live billing environments is needed to validate these findings.

The Layer 3 predictive model was trained on 2022–2023 claims data and reflects the payer behavior and policy landscape of that period. Behavioral health parity enforcement has been an area of active federal rulemaking and litigation; model performance may degrade over time as payer practices evolve, requiring regular retraining.

5.5 Future Directions

Several extensions of the BehaviorIQ framework merit investigation. First, prospective deployment in live billing environments with randomized design — BehaviorIQ versus general-purpose RCM AI for matched behavioral health laboratory claim populations — would provide the highest-quality evidence for the performance claims made in this retrospective evaluation. Second, extension of the MHPAEA parity module to include non-quantitative treatment limitation (NQTL) comparative analysis — documenting disparate step-therapy, prior authorization, and fail-first protocol requirements across behavioral health versus medical/surgical benefits — would expand the scope of parity-based appeal arguments available to the system. Third, integration of the 42 CFR Part 2 consent-conditional input gating architecture with patient-controlled consent management systems could enable real-time consent verification for individual claims, reducing the proportion of appeals that must use the restricted input pathway and potentially improving appeal success rates for the consent-absent claim population.

6. Conclusion

We present BehaviorIQ, a three-layer AI framework for behavioral health and addiction medicine laboratory claim compliance and denial recovery, built on the Deterministic-Generative-Predictive (DGP) Clinical Revenue Architecture. The framework addresses a documented and unmet need: behavioral health laboratory claims deny at rates approximately 58% above the all-specialties average, the existing AI literature contains no published work on AI-driven behavioral health laboratory claim compliance, and the unique regulatory requirements of 42 CFR Part 2 and MHPAEA create a compliance surface that general-purpose RCM AI systems cannot safely address without behavioral health-specific rule encoding.

The framework achieves a 17.6 percentage-point improvement in first-pass claim acceptance, a 74.8% appeal success rate, and 97.3% 42 CFR Part 2 compliance accuracy — materially exceeding both manual baseline and general-purpose AI comparator performance. The introduction of consent-conditional input gating as an architectural principle for 42 CFR Part 2-compliant AI systems represents a novel contribution to the healthcare AI literature.

BehaviorIQ establishes the foundational methodology for AI-assisted behavioral health laboratory revenue cycle management and introduces a research domain — AI for behavioral health claim compliance — that warrants continued investigation as the behavioral health parity enforcement landscape continues to evolve.

References

1. XiFin. *2024 Payor Denial Impact Report*. San Diego, CA: XiFin, Inc.; 2024.
 2. Kang SY, Odouard I, Gresenz CR. Claim Denials for Cancer-Related Next-Generation Sequencing in Medicare. *JAMA Netw Open*. 2025;8(4):e255785. doi:10.1001/jamanetworkopen.2025.5785.
 3. Kim BH, Sridharan S, Atwal A, Ganapathi V. Deep Claim: Payer Response Prediction from Claims Data with Deep Learning. arXiv preprint arXiv:2007.06229. 2020.
 4. Confidentiality of Substance Use Disorder Patient Records, 42 C.F.R. Part 2 (2024 revision). Available at: <https://www.ecfr.gov/current/title-42/chapter-I/subchapter-A/part-2>
 5. Substance Abuse and Mental Health Services Administration (SAMHSA). *TIP 63: Medications for Opioid Use Disorder*. SAMHSA Publication No. PEP21-02-01-002. Rockville, MD: SAMHSA; 2021.
 6. Mello MM, Trotsyuk AA, Djiberou Mahamadou AJ, Char D. The AI Arms Race In Health Insurance Utilization Review: Promises Of Efficiency And Risks Of Supercharged Flaws. *Health Affairs*. 2026;45(1):6-13. doi:10.1377/hlthaff.2025.00897. PMID 41494115.
 7. Johnson KB, Wei WQ, Weeraratne D, Frisse ME, Misulis K, Rhee K, Zhao J, Snowdon JL. Precision Medicine, AI, and the Future of Personalized Health Care. *Clin Transl Sci*. 2021;14(1):86-93. doi:10.1111/cts.12884. PMID 32961010; PMCID PMC7877825.
 8. Mental Health Parity and Addiction Equity Act of 2008, Pub. L. No. 110-343, § 512, 122 Stat. 3765 (2008).
-

Correction Note

Internal tracking only — not part of the submitted manuscript.

- **Ref 2** (formerly “Georgetown University Health Policy Institute... 2017–2021”, JAMA Netw Open): Replaced with the verified real article — Kang SY, Odouard I, Gresenz CR, JAMA Netw Open. 2025;8(4):e255785 — which studied Medicare claims data from 2016–2021, not 2017–2021. In-text reference to the study period (Section 2.1) corrected from “2017 to 2021” to “2016 to 2021” to match the real source; the incorrect “Georgetown University” institutional

attribution was also removed from the body text since it was not verifiable against the corrected citation.

- **Ref 3** (formerly “Kim D,” Deep Claim, arXiv:2007.06229): Corrected author listing to the verified full author set — Kim BH, Sridharan S, Atwal A, Ganapathi V — per fact-check findings.
- **Ref 6** (formerly a fabricated “Vadlamudi R, et al.” Health Affairs citation with a mismatched DOI): Removed entirely and replaced with the real, verified article — Mello MM, Trotsyuk AA, Djiberou Mahamadou AJ, Char D, Health Affairs. 2026;45(1):6-13 — which is a 2026 publication, not 2025 as originally stated; in-text year reference (Section 2.1) corrected from “Health Affairs 2025” to “Health Affairs 2026” to match.
- **Ref 7** (formerly “Johnson KB et al... Inf Syst Front 2023,” wrong journal/year): Replaced with the verified real citation — Johnson KB et al., “Precision Medicine, AI, and the Future of Personalized Health Care,” Clin Transl Sci. 2021;14(1):86-93. This real article is a general precision-medicine/AI review and does not report an NLP-based prior-authorization denial-prediction study with AUC 0.83 as the original body text claimed (Section 2.2); that fabricated, unsupported study description was removed and the surrounding paragraph lightly rephrased to preserve its valid point (that prior ML claims-denial literature is hospital-focused and does not address behavioral health parity or 42 CFR Part 2) without relying on invented findings. The corresponding in-text citation marker in Section 1.3 (“Existing RCM AI literature addresses hospital claims⁸...”) was also removed since it depended on the same unsupported claim.
- **All references renumbered** sequentially (1–8) in order of first appearance in the body text, and every in-text superscript citation marker was updated to match the new numbering. No reference was found to be “DELETE / fabricated, no real substitute” — Ref 1 (XiFin), Ref 4 (42 CFR Part 2), Ref 5 (SAMHSA TIP 63), and Ref 8 (MHPAEA 2008, formerly Ref 4) were verified accurate and retained without content changes, only renumbering.