



**ARDIA PRECISION HEALTH**

Next-Generation Intelligence for Healthcare Infrastructure


SEED ROUND · CONFIDENTIAL

**Governed, auditable AI  
for the healthcare work  
nobody automates.**

RAISING \$3-7M SEED

PRE-REVENUE

DALLAS-FORT WORTH · DE C-CORP · JAN 2026



## THE PROBLEM

# The crisis in lab reimbursement

# \$12B

in U.S. lab revenue is left unpaid every year — most of it appealable, simply written off. Industry est.



## The system is at a breaking point


~65% of denials are never appealed. Molecular & toxicology claims get denied as "experimental" despite guideline support — and a legislated PAMA rate cliff is about to make it worse.

## WHY NOW

# A legislated rate cliff hits labs in 2027–2029

- **PAMA / §1834A** lets CMS cut lab fee-schedule rates up to **15%/yr** across 2027, 2028, 2029.
- Compounding:  $0.85^3 = 0.614$  → a cap-bound code lands **~38.6% below** its 2025 rate by 2029.
- Failing to report private-payor data risks a CMP up to **\$10,000/day**.

THE PAIN IS SCHEDULED, NOT SPECULATIVE



## THE SOLUTION

# The governed reasoning + coverage layer for healthcare's hardest calls

### Recover denied revenue

Decode the denial, match payer policy + clinical guidelines, and draft a cited appeal — for toxicology & molecular labs.



### Model the PAMA cliff

**Meridian** — deterministic, unit-tested model of the 2027–2029 cuts on a lab's real test mix. Revenue-at-risk, per code.



### Enable precision care

Non-diagnostic decision-support for precision, pulmonary & elder care — and the coverage layer that gets tests paid.

CITE-OR-ABSTAIN

NON-DIAGNOSTIC

HUMAN-IN-THE-LOOP

NOTHING AUTO-FILED

THE PRODUCT · CORE

# TARA — the auditable reasoning core

A neuro-symbolic engine: a model reasons, a symbolic policy layer constrains it to real rules, and a guardrail layer must pass before any human sees the output.



## GROUNDING

De-identified data + live PubMed / ClinicalTrials.gov — no PHI reaches the model.

## SYMBOLIC POLICY

NCD / LCD / MoDX + deterministic PAMA math — the model can't invent a rule.

## GUARDRAILS

Crucible's 8 gates + cite-or-abstain — a draft for a licensed human, never auto-filed.

# TARA — the Triadic Adjudicative Reasoning Architecture

Neuro-symbolic: a language model reasons, a symbolic policy layer constrains it to real rules, and a guardrail layer must pass before any human sees the output.

## LAYER 1 · GROUNDING

### Data & De-identification

- EDI 835/837, HL7, FHIR ingestion
- Sentinel — HIPAA Safe-Harbor de-id
- Entity resolution & normalization
- Live PubMed / ClinicalTrials.gov
- Payer policy corpus (NCD/LCD/MoIDX)

NO PHI TO THE MODEL



## LAYER 2 · SYMBOLIC POLICY

### Rules the model can't break

- Coverage logic: NCD 90.2, MoIDX LCDs
- Coding: CPT/HCPCS, NCCI, MUE edits
- Deterministic math (Meridian / PAMA)
- Denial taxonomy (CARC/RARC)
- Every claim maps to a citable rule

CAN'T INVENT A RULE



## LAYER 3 · REASONING CORE

### Adjudicate, cite, guardrail

- LLM reasons over grounded evidence
- Swift · Cortex · Scholar tiers
- Crucible — 8 guardrail gates
- Cite-or-abstain, non-diagnostic
- Draft → human reviews & files

MUST PASS EVERY GATE

# One core, three depths — matched to the stakes

⚡ SWIFT · FAST

## Instant answer

Low-deliberation, no tool calls — a quick, direct response when speed matters.

🧠 CORTEX · THINKING

## Step-by-step

Shows its reasoning — decode → validate → find the rule → weigh evidence → draft.

🔍 SCHOLAR · DEEP RESEARCH

## Full investigation

Plans, calls tools, searches & retrieves policy + literature, then reasons and **cites**.

### EXAMPLE · SCHOLAR TRACE

**plan** → parse 835 (CARC 55) → map CPT 81455 → retrieve NCD 90.2 / MoIDX LCD → search NCCN + PubMed → reason → **guardrail check** → **cited**  
**appeal draft for a human**

HOW IT WORKS

# From a denied claim to a cited appeal



## AUDITABLE BY DESIGN

Every step logged; every claim traces to a cited policy or study.

## DE-IDENTIFIED END-TO-END

PHI stripped before any model call — even on the emergency path.

## HUMAN-IN-THE-LOOP

Output is a draft; a licensed person verifies and submits.



THE MODELS

# Nine models & kernels — one governed core



✓ **MEASURED** real, tested code · **IN BUILD** persona on a frontier model under guardrails, no published accuracy yet

# A deterministic model of the PAMA cliff

- Enter a lab's real test mix; Meridian compounds the 15%/yr cuts **year by year** → revenue-at-risk **per code**.
- Deterministic — **no LLM, no PHI**. The math is **unit-tested** & reproducible.
- Seed mix: baseline **\$895,551** → 3-yr shortfall **\$728,419**. *illustrative*

## THE MATH

$0.85^3 = 0.6141 \rightarrow 38.6\%$  cut from 2025 by 2029



## TRUST & SAFETY

# In healthcare, traceability **is** the product

### **SENTINEL — DE-IDENTIFICATION**

18 HIPAA Safe-Harbor identifiers stripped before any model call or log write.  
Encryption at rest / in transit, hash-chained audit, minimum-necessary access.

### **CRUCIBLE — 8 GUARDRAIL GATES**

Non-diagnostic · cite-or-abstain · safety-escalation · scope-of-practice · honesty  
· human-in-the-loop. An output must pass every applicable gate before a person  
sees it.



HIPAA SAFE HARBOR

BAA-GATED BEFORE REAL PHI

TX SB 1188 · TRAIGA AWARE

SOC 2 · GCP — ROADMAP

# Pre-revenue — but real, tested code

95.45%

Cadence held-out accuracy · public UCI HAR,  
subject-independent, reproducible

34/34

automated tests green in CI · Cadence ·  
Meridian · Sentinel · Crucible

Live

PubMed + ClinicalTrials.gov grounding with  
cite-or-abstain in the Studio

## The honest line.

One model (**Cadence**) is genuinely trained and measured; **Meridian**'s math is unit-tested; the persona models run on a frontier model under guardrails with **no published accuracy yet** — benchmarks will be built the same way, on real data under a BAA. A live Studio runs the pipeline end-to-end today. **Nothing invented; none of it a diagnosis.**

## MARKET

# A large, non-discretionary problem

### TAM • U.S. clinical lab economy

\$100B+

Clinical lab testing + revenue-cycle management. industry est.

### SAM • Denial recovery + PAMA exposure

\$ billions

Molecular/tox denials never appealed + labs modelling the cliff.

### SOM • Beachhead

DFW → regional

Design-partner independent labs in Dallas-Fort Worth, then the Sun Belt.



## BUSINESS MODEL

# We get paid when the lab gets paid

### PRIMARY

#### Contingency on recovery

A % of revenue Ardia helps recover on wrongly-denied claims. Aligned incentives, near-zero adoption risk.

### EXPANSION

#### Platform SaaS

Subscription for the Studio, Meridian PAMA modelling, and coverage intelligence across a lab's book.

### FUTURE

#### Precision-care layer

Decision-support + coverage enablement for precision, pulmonary & elder care.

**WHY IT LANDS** **No budget fight** (contingency comes out of newly-recovered dollars) · **fast proof** (recover on a denial backlog in weeks) · **sticky** (the audit trail + PAMA model become operational infrastructure). Unit economics modelled; validated with design partners.

## ROADMAP

# Timed to the cliff

2026

### Build & pilot

Reasoning core, Meridian, Sentinel/Crucible tested. DFW design-partner labs. First recoveries. SOC 2 & BAA underway.

2027

### Cliff year 1

PAMA cuts begin. Recovery + Meridian live. Expand tox & molecular coverage. Grow the book.

2028–29

### Scale

Regional expansion across the Sun Belt. Pulmonary & specialty lines. Payer-facing evidence.


2030–32+

### Precision platform

Precision, pulmonary & elder care decision-support at scale on the governed core.

## COMPETITION

# Others automate the easy middle. We own the auditable tail.



### Our wedge

- **Auditable reasoning** — every claim cites a policy or study
- **Long-tail focus** — the molecular/tox denials bots can't touch
- **Governed** — de-id, guardrails, human-in-the-loop
- **PAMA-native** — Meridian is a reason to buy *now*

### Moat

The audit trail + payer-policy corpus + design-partner data compound into switching costs and coverage intelligence.



## TEAM

# Built by operators who ship real, tested code



### Ram Vadlamudi

Founder · Creator of TARA

Architect of Ardia's Triadic Adjudicative Reasoning Architecture — the governed, auditable core the platform runs on.



### Paramesh Kurapati

Chief Executive Officer

Leads company strategy, partnerships, and go-to-market.



### Manasa Jampani

Co-founder

Co-founding team — product, delivery, and operations.

**Already shipped:** a live platform ([ardiahealthlabs.com](https://ardiahealthlabs.com)) with a working Studio, the tested Meridian / Sentinel / Crucible / Cadence components, 34 tests in CI, and live PubMed/ClinicalTrials grounding — evidence this team executes.

## THE ASK

# Raising a \$3–7M seed

to reach first recoveries & design-partner revenue ahead of the 2027 cliff.

### USE OF FUNDS

- Engineering — productionize the reasoning core & recovery pipeline
- Clinical & regulatory — BAA, SOC 2, payer-policy corpus
- Go-to-market — DFW design-partner labs & first recoveries

### 18-MONTH MILESTONES

- Signed design-partner labs under BAA
- First measured revenue recovered on real denials
- SOC 2 Type I & production security
- Published model benchmarks (the Cadence way)

**\$3–7M**

Seed round

**18mo**

Runway to revenue  
proof

**2027**

Cliff = forcing  
function

**DEC-Corp**

Clean cap table



**ARDIA PRECISION HEALTH**

Next-Generation Intelligence for Healthcare Infrastructure

**The right care, delivered and  
paid for — every decision  
auditable.**

[ARDIAHEALTHLABS.COM](https://ARDIAHEALTHLABS.COM)

[INFO@ARDIAHEALTHLABS.COM](mailto:INFO@ARDIAHEALTHLABS.COM)

SEED · \$3–7M